

Computer Abstractions and Technology

CSI3102-02: Architecture of Computers
(컴퓨터아키텍처)

Youngsok Kim (김영석)



Computer Architecture

- A set of rules and methods which describe the functionality, organization, and implementation of **computer systems**
- In other words, **computer architects** design the internals of computer systems!
- Computer systems
 - A complete **computer** including the hardware, the operating system, and peripheral equipment required for performing full operations



Computers

- Machines which can be instructed to **execute series of arithmetic or logic operations automatically** via **computer programming**
- Via computer programming
 - Perl/Python/Ruby, C/C++, Java, assembly, ...
- Execute arithmetic or logic operations
 - Instructions generated by compilers/interpreters
- Execute the operations **automatically**
 - This is what you will learn throughout this semester!



Traditional Classes of Computers

- Personal Computers (PCs)
 - Good performance to single users at low cost
- Servers
 - Greater computing, storage, and input/output capacity
- Supercomputers
 - The high-end extreme of servers
- Embedded Computers
 - Run one application or one set of related applications typically integrated with the hardware


Q: Why split computers to the classes in the first place?

A: Because the performance/cost requirements differ!



The Post-PC Era

- **Personal Mobile Devices (PMDs)** instead of PCs
 - Operated with a battery and wireless connections
 - Likely to rely on a touch-sensitive screen or speech input



The manufactured # of PMDs (tablets, smart phones) now exceed PCs and traditional cell phones, reflecting the post-PC era.

- **Cloud Computing** instead of servers
 - a.k.a. Warehouse Scale Computers (WSCs)
 - Provide Software-as-a-Service (SaaS) via the cloud
 - e.g., Google App Engine, Amazon Web Services, Microsoft Azure



Below Your Program



Abstracting a Computer System

- **Application software**


- High-level software running on top of the systems software
- e.g., web browsers, games

- **Systems software**

- Low-level software for executing the application software on the hardware
- e.g., operating systems, compilers

- **Hardware**

- The underlying hardware that executes operations from the systems software
- e.g., CPUs, GPUs



Operating Systems & Compilers

- Two types of systems software tightly coupled to the underlying hardware
- **Operating Systems (OSes)**
 - Interfaces between user applications and the hardware
 - Handle basic I/O operations, allocate memory/storage, ...
 - e.g., Windows, Linux, OS X
- **Compilers**
 - Translates a program written in a high-level language into instructions which the underlying hardware can execute
 - e.g., gcc/g++, LLVM



Compiling a User Program

- **Assembly** language
 - Utilizes symbolic notations to abstract the machine language
 - e.g., add A, B instead of 1001010100101110
- **Binary machine** language
 - Only consists of bits (0s & 1s)
 - **THE** language electronic hardware can understand

High-level
language
program
(in C)

```
swap(size_t v[], size_t k)
{
    size_t temp;
    temp = v[k];
    v[k] = v[k+1];
    v[k+1] = temp;
}
```

Compiler

Assembly
language
program
(for RISC-V)

```
swap:
    slli x6, x11, 3
    add x6, x10, x6
    ld x5, 0(x6)
    ld x7, 8(x6)
    sd x7, 0(x6)
    sd x5, 8(x6)
    jalr x0, 0(x1)
```

Assembler

Binary machine
language
program
(for RISC-V)

```
00000000001101011001001100010011
00000000011001010000001100110011
00000000000000110011001010000011
00000000100000110011001110000011
00000000011100110011000000100011
00000000010100110011010000100011
00000000000000001000000011001111
```



Under the Covers & Building Processors and Memory



Components of a Computer

- **Datapath & Control** (the Processor)
 - Get instructions and data from memory, and execute the instructions using the data
- **Memory**
 - Serves as an intermediate storage for instructions and data between the processor and the I/O
 - Organized in a hierarchical fashion
 - e.g., two-level memory hierarchy with SRAM and DRAM
- **Input & Output (I/O)**
 - Write data to memory & read data from memory
 - e.g., keyboards, mouse, monitors, speakers



Inside Apple iPad 2

- Contains all the 5 components of a computer
 - Processor: Apple A5
 - Memory: 512 MiB
 - I/O: touchscreen, LCD, WiFi, ...
- Integrated Circuits (Chips)
 - 2x 1-GHz ARM processors
 - 2x 2-Gib DRAM chips



Components of Apple iPad 2



The logic board of Apple iPad 2

Apple iPad 2's Microprocessor

- Datapath
 - Perform arithmetic operations
 - e.g., two ARM processors, a PowerVR graphics processing unit
- Control
 - Instructs the datapath, memory, and I/O devices
 - e.g., interfaces to DRAM, I/O, WiFi, Audio



The Hardware/Software Interface

- **Instruction Set Architectures (ISAs)**

- The interface btw. the H/W and the lowest-level S/W
- Include anything programmers need to know to make a binary machine language program work correctly
 - e.g., instructions, I/O devices

- **Application Binary Interfaces (ABIs)**

- The basic instruction set & the OS interface provided for application programmers
- Encapsulate the details of I/O, memory allocation, etc. (by the underlying OS)



How Do We Store Data?

- **Volatile** memory

- Forget the data when it loses power
- Used for holding data and programs inside the computer
 - e.g., Dynamic Random Access Memory (DRAM) chips
- Fast, but small and expensive
- a.k.a. **main/primary** memory

- **Non-volatile** memory

- Maintain the data as is even though it loses power
 - e.g., Hard Disk Drives (HDDs), Solid-State Drives (SSDs)
- Large and cheap, but slow
- a.k.a. **secondary** memory

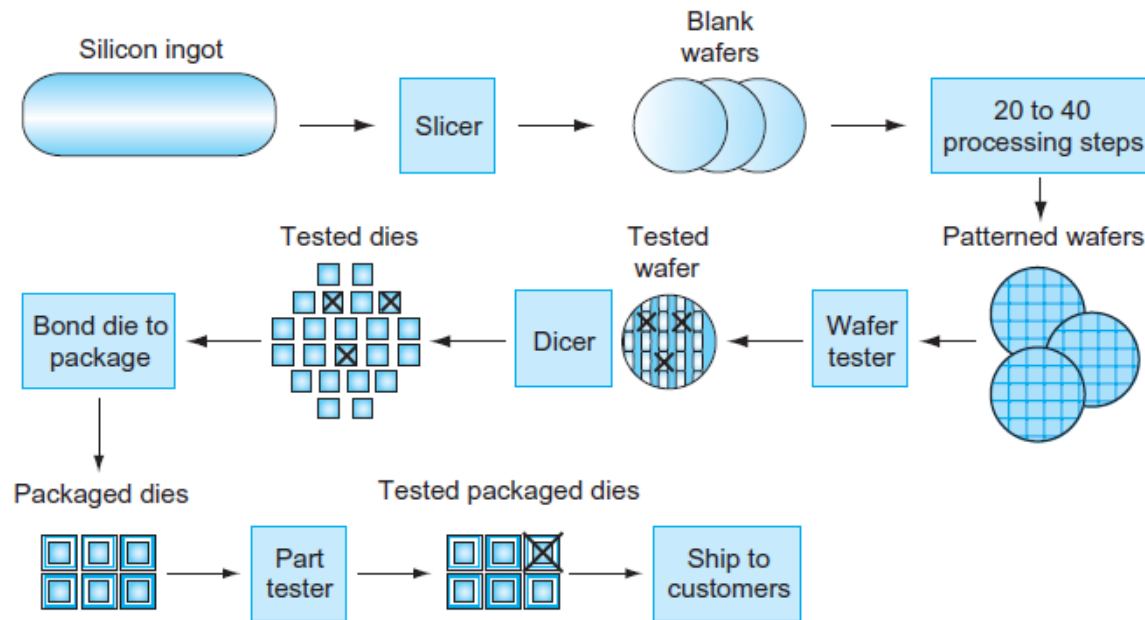


Manufacturing Integrated Circuits

- An Integrated Circuit (IC) consists of **transistors**.
 - On/off switches controlled by electricity

- Silicon-to-IC

Silicon →
Blank wafers →
Patterned wafers →
Tested dies →
Packaged dies →
Tested packaged
dies (or **chips**)



Performance



Our Focus: “Make Programs Fast”

- Let’s think of ourselves as **successful** programmers who are always concerned about the **performance** of our applications!
- To improve the performance, we must be aware of the primary performance **bottlenecks**.
 - In 1960s & 70s, it was the size of a computer’s memory.
→ minimize memory space to make programs fast!
 - Now, the bottlenecks are the **parallel nature of processors** & the **hierarchical nature of memories**!
 - On PMDs and in the Cloud, we also need to consider **energy efficiency**!



Defining Performance

- When we say one computer is **faster** than another, what does it exactly mean?
- We need **performance metrics!**
 - 1. Execution time**
 - The total time required for a computer to complete a task
 - Useful for single-user desktops
 - 2. Throughput** (or bandwidth)
 - The total amount of work done in a given time
 - Useful for servers or WSCs which run various tasks concurrently

Comparing the Performance

- Let's say fast execution time == high performance.

$$\text{Performance}_X = \frac{1}{\text{Execution Time}_X}$$

- Now, we compare the performance of two computers:

$$\text{Performance}_X > \text{Performance}_Y$$

$$\frac{1}{\text{Execution Time}_X} > \frac{1}{\text{Execution Time}_Y}$$

$$\text{Execution Time}_Y > \text{Execution Time}_X$$

→ X is faster than Y ==

The execution time on Y is longer than that on X.



Quantifying the Performance

- We say “X is **n times faster than Y.**” when:

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = n$$

- That is,

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = \frac{\text{Execution Time}_Y}{\text{Execution Time}_X} = n$$

- In the other direction, “Y is **n times slower than X.**”

$$\frac{\text{Performance}_X}{\text{Performance}_Y} = n \qquad \text{Performance}_Y = \frac{\text{Performance}_X}{n}$$

Measuring the CPU Performance

- Is using the **wall clock time** the correct measure?
- Multiple tasks run in parallel even on single-user PCs.
→ Need to distinguish between the elapsed time & the time the CPU spent working on our tasks!
- **CPU execution time (or CPU time)**
 - The time the CPU spends computing for our tasks excluding time spent waiting for I/O or other tasks
 - Can be further divided into **user** & **system** CPU times
 - User CPU time: the CPU time spent in the program
 - System CPU time: the CPU time spent in the OS



Using Clock Cycles as the Measure

- Processors are **digital** circuits driven by a clock!
 - All events in a processor occur w.r.t. the clock signal.
- Use **the length of a clock period** as the measure
 - The time for a complete clock cycle
 - e.g., 1 ns for a 1-GHz clock, 250 ps for a 4-GHz clock
- Then, we get:

$$\begin{aligned}\text{CPU execution time} &= \text{CPU clock cycles} \times \text{Clock cycle time} \\ &= \text{CPU clock cycles} / \text{Clock rate}\end{aligned}$$



Example #1

- Our program runs in 10 seconds on computer A having a 2-GHz clock.
- Q: How many clock cycles does the program take?

$$\text{CPU time}_A = \text{CPU clock cycles}_A / \text{Clock rate}_A$$

$$10 \text{ seconds} = \text{CPU clock cycles}_A / 2 \times 10^9 \frac{\text{cycles}}{\text{second}}$$

$$\text{CPU clock cycles}_A = 10 \text{ seconds} \times 2 \times 10^9 \frac{\text{cycles}}{\text{second}} = 20 \times 10^9 \text{cycles}$$



Example #2

- On computer B, our program takes 6 seconds, and 1.2 times as many clock cycles as computer A.
- Q: What is the clock rate/frequency of computer B?

$$\text{CPU time}_B = \frac{\text{CPU clock cycles}_B}{\text{Clock rate}_B}$$

$$6 \text{ seconds} = \frac{1.2 \times \text{CPU clock cycles}_A}{\text{Clock rate}_B}$$

$$6 \text{ seconds} = \frac{1.2 \times 20 \times 10^9 \text{ cycles}}{\text{Clock rate}_B}$$

$$\text{Clock rate}_B = \frac{24 \times 10^9 \text{ cycles}}{6 \text{ seconds}} = 4 \text{ GHz}$$



Adding Instructions to Performance

- A program consists of **instructions** generated by compilers and interpreters.
- Use **Cycles Per Instruction (CPI)** instead:

$$\text{CPU clock cycles} = \text{Instruction Count} \times \text{Average clock cycles per instruction}$$

- Instructions may take different amounts of time, so we use average CPI.

The Classic Performance Equation

- Using the CPI, the CPU time becomes:

$$\begin{aligned}\text{CPU time} &= \text{CPU clock cycles} / \text{Clock rate} \\ &= \text{Instruction count} \times \text{CPI} / \text{Clock rate} \\ &= \text{Instruction count} \times \text{CPI} \times \text{Clock cycle time}\end{aligned}$$

- Alternatively, we can write:

$$\begin{aligned}\text{Time} &= \text{Seconds} / \text{Program} \\ &= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}}\end{aligned}$$

Eight Great Ideas in Computer Architecture



The Eight Great Ideas (1/2)

- **Moore's Law**

- "Integrated circuit resources double every 18-24 months."

- **Abstraction**

- Abstract/characterize the hardware design at different levels of representations

- **Make the Common Case Fast** (Amdahl's Law)

- e.g., If functions A and B take 80% and 20% of the time, focus on function A rather than function B.

- **Parallelism**

- Improve the performance of applications by executing operations in parallel



The Eight Great Ideas (2/2)

- **Pipelining**

- e.g., divide laundry into three steps (e.g., wash, dry, fold), and fold clothes A while drying/washing clothes B/C

- **Prediction**

- Guess & start executing next instructions instead of waiting

- **Hierarchy of Memories**

- Place a faster, smaller, expensive memory near processors, and a slower, larger, and cheaper memory farther

- **Dependability via Redundancy**

- Redundantly execute operations on redundant hardware, and compare the outcomes from the redundant hardware



Make the Common Case Fast

- “Improving one aspect of a computer would increase the overall performance of the computer!”
 - e.g., Improving A of a CPU by 20% improves the performance of the CPU by 20%.

- **Amdahl's Law**

Execution time after improvement =
$$\frac{\text{Execution time affected by improvement}}{\text{Amount of improvement}} + \text{Execution time unaffected}$$

- e.g., Making 20% of a 5-sec. application 1.2 times faster:

$$\text{Time} = \frac{0.2 \times 5}{1.2} + (0.8 \times 5) = 4.83 \quad \leftarrow \text{only 0.17 seconds faster!}$$




The Power Wall & The Switch to Multiprocessors



Trends in Clock Rate and Power

- Clock rate and power increased rapidly for decades. Then, they flattened off recently.



We ran into the practical power limit for **cooling!**

Clock Rate – Power Relationship

- ICs are typically made with **CMOS**.
 - Complementary Metal-Oxide Semiconductor (CMOS)
- In CMOS, **dynamic energy** is the primary source of energy consumption!
 - Depends on the **capacitive loading** of each transistor and the **supply voltage**:

$$\text{Energy} \propto \text{Capacitive load} \times \text{Voltage}^2$$

- The energy of a single transition (0→1 and 1→0) is:

$$\text{Energy} \propto \frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2$$

- Then, the power required per transistor becomes:


$$\text{Power} \propto \frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency switched}$$

where *Frequency switched* is a function of **the clock rate**.



The Switch to Multiprocessors

- Can't increase the clock rate due to the power wall
- Okay, let's add more **cores** instead!



The Era of Parallel Programming

- The free lunch is now over :(
 - Your program doesn't get faster anymore as the CPU's clock rate doesn't increase.
- Need to exploit **parallelism** for high performance; however, it has been & still is difficult.
 - Parallel programming **increases programming difficulty**.
 - Programmers **must divide an application** so that each processor has roughly the same amount of work.
 - Programmers **must minimize inter-processor communication and synchronization**.

